

# Missing Data - Assignment 1

Aga, Nisse, Ruben

2024-02-28

## Contents

<b>1</b>	<b>Introduction (Ruben)</b>	<b>1</b>
<b>2</b>	<b>Methodology (Aga)</b>	<b>2</b>
2.1	Dataset . . . . .	2
2.2	Data variables . . . . .	2
2.3	Data processing methodology . . . . .	3
<b>3</b>	<b>EDA Results (Nisse)</b>	<b>4</b>
3.1	Descriptive statistics . . . . .	4
3.2	Correlations . . . . .	6
3.3	Outliers . . . . .	7
<b>4</b>	<b>Missing data problem (Aga)</b>	<b>7</b>
4.1	Missing data and response Patterns . . . . .	7
4.2	Model creation (maybe combine with next section?) . . . . .	12
4.3	Comparison of the two different models in terms of missing data treatment !!! (Ruben) . . . .	14
4.4	Conclusion in terms of answering RQ (Nisse) . . . . .	14
<b>5</b>	<b>References</b>	<b>14</b>
<b>A</b>	<b>Appendix</b>	<b>15</b>

## 1 Introduction (Ruben)

### 1.0.1 RQ (Ruben)

Moore et al. (2005) Garnett et al. (2022) Wu, Jia, and Enders (2015) Alcohol consumption led to 2.8 million deaths in 2016 and accounts for almost 10% of global deaths in people aged 15-49 years. Higher levels of alcohol consumption leads to a higher risk of mortality and the only level of alcohol consumption that minimizes this risk is zero alcohol consumption (study). This means that drinking any alcohol at all

increases the risk of mortality. All of this makes it vital to know and understand the factors behind alcohol consumption.

Moore et al. 2005 found that age, sex, ethnicity, marital status, education level and household income were all either negatively or positively associated with alcohol consumption. In Garnett et al. 2022 it was found that the level of depression was negatively associated with alcohol consumption. Considering that the only safe level of alcohol consumption is zero, it is important to understand what factors predict regular consumption of alcohol instead of only looking at the amount of alcohol consumption as in the two studies mentioned above.

The research question of this study is whether the occurrence of regular alcohol consumption (12 or more in a year) can be predicted by the variables: depression level, age, sex, ethnicity (white vs other), marital status and household income. It is expected that age and depression level will be negatively correlated with the occurrence of alcohol consumption while household income will be positively correlated (sources). It is also expected that being male (vs female) and white (vs other ethnicities) will be positively correlated with the regular occurrence of alcohol consumption (source). Regarding marital status it is expected that the married status will be positively associated with the regular occurrence of alcohol consumption compared to others but that is based on a study where the factor marital status consisted of just married and other (source), while in this study more categories are considered.

## 2 Methodology (Aga)

### 2.1 Dataset

The dataset used is a subset of the data collected in the National Health and Nutrition Examination Survey (NHANES). The survey is a part of annual program that investigates the health and nutrition of a representative sample of people in the United States. The data we used contains information about 525 individuals that has been collected for the NHANES 2007-2008. This is a subset of the 12,946 individuals in that years' survey sample, out of which 78.4% was interviewed and 75.4% was examined in mobile examination centers. The NHANES is further subdivided into themed sections, such as the Alcohol Use questionnaire, that have separate documentation that will be referred to later.

The used dataset contains a wide range of variables related, to the health of the individuals. We further subset the data by only including variables relevant to the study (demographics, alcohol use and answers to depression screener questions). The selected variables are further described in Variable Description section.

### 2.2 Data variables

Table 1 lists the variables used in our subset selection, which will be utilised for the model in question. The predictor variables [*dep1...dep9*] are sourced from the same Depression Screener, where respondents of age 18 to 150 were ought to assign a number (1 to 3) regarding their mental and physical state within the last 2 weeks. Multiple signs of depression were measured this way, which can later be combined to indicate an overall level of depression.

The demographic variables - that being **sex**, **age**, **ethnicity**, **education** and **household\_income** - were taken from the same screening component as well. The following should be noted, regarding these demographic variables:

- The variable **age** is topcoded at the value 80 for the respondents who were older than 80 years.
- The variable **education** was targeted at respondents of age 20 to 150, thus excluding younger participants. This is due to the fact that this question includes responses such as **AA degree** and **College Graduate**.
- Similarly, the variable **marital** was also targeted at respondents of age 20 to 150.

Table 1: Variable descriptions

Role	Variable	Name	Type	Characteristics	Target
Outcome	Drink regularly	drink_regularly	Categorical	Binary, yes and no	m/f, age 20-150
Predictor	Sex	sex	Categorical	Binary, male and female	m/f, age 0-150
Predictor	Age	age	Numeric	Discrete	m/f, age 0-150
Predictor	Ethnicity	ethnicity	Categorical	Nominal, 5 categories	m/f, age 0-150
Predictor	Education	marital	Categorical	Nominal, 5 categories	m/f, age 20-150
Predictor	Marital status	marital	Categorical	Nominal, 5 categories	m/f, age 20-150
Predictor	Household income	household_income	Categorical	Nominal, 12 categories	m/f, age 0-150
Predictor	No interest in activity	dep1	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Feeling depressed	dep2	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Sleeping issues	dep3	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Feeling tired	dep4	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Eating issues	dep5	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Feeling bad about yourself	dep6	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Concentrating issues	dep7	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Moving and speaking issues	dep8	Categorical	Ordinal, 1-3 scale	m/f, age 18-150
Predictor	Suicidal thoughts	dep9	Categorical	Ordinal, 1-3 scale	m/f, age 18-150

- The variable `household_income` is ordinal, rather than continuous.

As for the remaining demographic variables, namely `sex`, `age`, `ethnicity` and `household_income`, these are retrieved from target age 0 to 150.

Finally, the `drink_regularly` variable was obtained from an Alcohol Use questionnaire targeted at ages 20 and up.

## 2.3 Data processing methodology

### 2.3.1 Data processing methodology

After arriving at the subset dataset, as Exploratory Data Analysis was performed. The distributions of the variables were investigated and presence of missing data was found in the outcome variable and some of the depression questions. The important findings of this step are presented in ‘EDA Results’ section.

Following the EDA, the missing data problem was addressed. The extend, distribution and patterns of missingness were investigated. Testing was done to verify the missing data mechanism, including testing of the dependencies between missing values in one variable and observed values of other variables and performing the Little (1988) MCAR Test.

Two solutions to the missing data problem were implemented: list-wise deletion and mean imputation. These were chosen based on the convenience of implementation. For list-wise deletion all cases with one or more missing values were excluded from the modeling. The mean imputation method was implemented in the following way: For the missing depression questions the mean of the observed values within the variable was computed and then the missing values were replaced by the mean value. For ‘`drink_regularly`’, the outcome variable, instead of a mean, the mode was computed. This was motivated by the need for a binary outcome variable in the logistic regression that was planned to answer the research question. Since the EDA found there was significantly more ‘yes’ values than ‘no’ values in ‘`drink_regularly`’, the missing values were replaced by ‘yes’. Thus, two complete datasets with different missing data treatments were created.

### 2.3.2 Model methodology

Including all of the variables in the two complete data sets is motivated by theoretical findings since other researchers found a relationship between them and drinking habits. Therefore, verifying the significance of these predictors and their relative importance in the presence of other variables was important. Thus, a decision was made to create two logistical models including all of the variables, instead of a top down or bottom up approach to model building. As opposed to removing the non-significant predictors as in the other approaches, we decided to keep them and therefore have more complex models.

The two logistics models were then compared and the impact of the two different missing data treatments was evaluated.

- what we investigated and why

## 3 EDA Results (Nisse)

### 3.1 Descriptive statistics

Table 4 from the appendix shows summary statistics for each of the variables within the dataset; including mean values, standard deviations, IQR statistics and data range values. For categorical variables, a list of possible categories and their respective proportions is provided to replace the continuous summary statistics. Lastly, the column N shows the amount of cases with present data - that being non-NA values.

In general, it is worthy to note that all variables are interpreted as a factor, excluding **age** and the multiple depression levels of **dep**. Table 1 suggests that **dep** should in fact be categorical ordinal and should therefore be cast as a factor. That said - and as mentioned in Section ?? - keeping the levels of **dep** as a numerical continuous datatype is beneficial for our missing data problem.

Our dichotomous outcome variable **drink\_regularly** has 307 cases of “yes” and 107 cases of “no”, having a outcome balance of 69% and 31% respectively. This outcome ratio could be considered imbalanced, which could affect the accuracy of our logistic regression model. Additionally, the total amount of value entries (N) of 446 suggests that 79 cases contain values outside of the set of possible binary values - most likely being missing values. Table 2 further confirms this.

Table 2: Drink regularly value distributions

drink_regularly	n
yes	307
no	139
NA	79

Predictor **sex** is a dichotomous variable with a balanced frequency distribution across the values “male” and “female”, that being 48% and 52% respectively. 525 entries contain one of these values, suggesting that the variable does not contain missing data.

The predictor **age** is the only continuous variable present within the sub-selected dataset. Although the survey was targeted at respondents of age 0-150 for most variables - the documentation even mentioning the topcoded entries for age 80+ - the dataset seems to only contain cases of people between the age of 20 and 69. Moreover, Figure 1 suggests a uniform distribution of the age variable. Like **sex**, **age** has 525 cases with non-NA values, hence the variable does not contain missing data values.

Predictor **ethnicity** has 5 categories, with **non-hispanic\_white** being the most prevalent category with 220 cases (42% of total), all the while **other** is the most infrequent category with 25 cases (5%). A total




Figure 1: Distribution of age

of 525 cases contain **ethnicity** data, once again suggesting no presence of missing data values within this variable.

Predictor **education** is similar to **ethnicity**, having 5 categories. That said, the frequencies of said categories seem to be less out of proportion, with **some\_college** being the most frequent category with 155 (30%) cases. Whilst the data was retrieved from respondents of ages 20 and higher, we do not observe any missing data values within the variable. This can be further justified by the fact that the minimum **age** within our dataset is 20, therefore foregoing possible issues with younger participants being unable to provide data for this specific variable.

The predictor variable **marital** contains 6 categories. Prior to combining the **marital** categories, the category value **married** exceeded the the average frequency of other categories (that being 49.2) by a large margin - as can be seen in Figure 2. Specifically, 279 cases (53%) were **married**, with the other 5 categories made up the remaining 47% of the cases. **never\_married** was the most frequent among those other 5 categories with 102 (19%) cases, whilst **separated** was the most infrequent category with only 14 (3%) cases.




Figure 2: Distribution of marital status categories before pre-processing

### 3.2 Correlations

```
## Warning in chisq.test(data$marital, data$drink_regularly): Chi-squared  
## approximation may be incorrect
```

```
##
```

Table 3: Contingency table of marital vs. drink regularly

	yes	no
married	175	81
widowed	8	11
divorced	48	14
separated	8	4
never_married	45	23
living_with_partner	23	6

```
## Pearson's Chi-squared test
##
## data: data$marital and data$drink_regularly
## X-squared = 10.218, df = 5, p-value = 0.06928
```

```
##              X^2 df P(> X^2)
## Likelihood Ratio  9.8952  5 0.078259
## Pearson          10.2183  5 0.069280
##
## Phi-Coefficient   : NA
## Contingency Coeff.: 0.15
## Cramer's V       : 0.151
```

TODO (from pract):

- Variable with most NA's
- Summary of non-NA vars
- 

### 3.3 Outliers

For our only continuous predictor **age**, a boxplot shows potential outliers if datapoints fall outside of the interquartile range (whiskers of the plot). Figure 3 shows the distribution of **age** in said boxplot, revealing no possible outliers. This is to be expected, since **age** was uniformly distributed as mentioned in Section 3.1.

## 4 Missing data problem (Aga)

### 4.1 Missing data and response Patterns

Firstly, we investigate the overall distribution of missing data in our dataset:






Figure 3: Boxplot of age





As can be seen on the graph above, 8.2% of the data is missing. The missing values occur in the outcome variable 'drink\_regularly' and in the responses to questions 'dep2', 'dep3', 'dep5', 'dep6', 'dep8' and 'dep9' that create the depression score variable. 15% of responses are missing for the predictor variable. 25% of the responses are missing for the individual depression questions 2, 3, 5 and 6, 10% are missing for question 8 and 14% for question 9.

We further investigate the missing data patterns by looking at the response patters:



##	id	sex	age	ethnicity	education	marital	household_income	dep1	dep4	dep7	dep8
##	263	1	1	1	1	1		1	1	1	1
##	72	1	1	1	1	1		1	1	1	1
##	45	1	1	1	1	1		1	1	1	1
##	22	1	1	1	1	1		1	1	1	1
##	43	1	1	1	1	1		1	1	1	1
##	19	1	1	1	1	1		1	1	1	1
##	7	1	1	1	1	1		1	1	1	1
##	2	1	1	1	1	1		1	1	1	1
##	29	1	1	1	1	1		1	1	1	0
##	15	1	1	1	1	1		1	1	1	0
##	2	1	1	1	1	1		1	1	1	0
##	1	1	1	1	1	1		1	1	1	0
##	5	1	1	1	1	1		1	1	1	0
##		0	0	0	0	0		0	0	0	52
##	dep9	drink_regularly	dep2	dep3	dep5	dep6					
##	263	1	1	1	1	1	0				
##	72	1	1	0	0	0	4				
##	45	1	0	1	1	1	1				
##	22	1	0	0	0	0	5				
##	43	0	1	1	1	1	1				
##	19	0	1	0	0	0	5				
##	7	0	0	1	1	1	2				
##	2	0	0	0	0	0	6				
##	29	1	1	1	1	1	1				
##	15	1	1	0	0	0	5				

```
## 2      1      0      1      1      1      1      2
## 1      1      0      0      0      0      0      6
## 5      0      1      1      1      1      1      2
##      76      79     131     131     131     131    731
```

This figure reveals that there are 13 distinct response patterns in the dataset and the missingness pattern is not monotone. The most frequent pattern is no missing entries, with 263 cases. It is important to note that the depression questions 2, 3, 5 and 6 are always either all present or all missing. It is very probable that the reason for item non-response for the depression items is the same, since there are no cases of only some of them missing.

Based on the missingness pattern of the depression items, 41% of the overall depression score includes at least one missing value.

#### 4.1.0.1 Testing dependency of missing values

```
##      t
## 19.31658

## [1] 3.099076e-45

## # A tibble: 1 x 4
##   statistic    df p.value missing.patterns
##   <dbl> <dbl> <dbl>         <int>
## 1    471.   164      0             13
```

Thus, the missing values are definitely not missing at random.

- what's the missing data mechanism?

#### 4.1.1 Result models with deletion and imputation (Nisse)

- formula
- table with coefficients and pval (make sure to exponential the coefficients for easier interpretation)
- Interpretation of model result

```
miceOut <- mice(data, defaultMethod = c("norm.predict", "logreg", "polyreg", "polr"), m = 1, maxit = 1)
```

```
##
## iter imp variable
## 1 1 drink_regularly dep2 dep3 dep5 dep6 dep8 dep9
```

```
reg_imp_data <- complete(miceOut)
summary(reg_imp_data)
```

```
##      id      drink_regularly      sex      age
## Min.   :41531  yes:364      male :254  Min.   :20.00
## 1st Qu.:43912  no :161      female:271  1st Qu.:33.00
## Median :46357                                Median :45.00
## Mean   :46470                                Mean   :44.99
```

```

## 3rd Qu.:48934                      3rd Qu.:57.00
## Max.      :51610                    Max.      :69.00
##
##          ethnicity          education          marital
## mexican_american : 95  no_high_school : 58  married          :279
## other_hispanic    : 61  some_high_school:101  widowed           : 19
## non-hispanic_white:220  high_school_grad:123  divorced          : 67
## non-hispanic_black:124  some_college    :155  separated         : 14
## other              : 25  college_grad   : 88  never_married     :102
##                                  living_with_partner: 44
##
##      household_income      dep1      dep2      dep3
## 100000+      : 76      Min.      :0.0000  Min.      : -0.3146  Min.      : -0.24112
## 25000:34999 : 59      1st Qu.:0.0000  1st Qu.: 0.0000  1st Qu.: 0.00000
## 20000:24999 : 52      Median :0.0000  Median : 0.0000  Median : 0.08312
## 35000:44999 : 51      Mean      :0.4095  Mean      : 0.3664  Mean      : 0.69010
## 75000:99999 : 49      3rd Qu.:1.0000  3rd Qu.: 0.7531  3rd Qu.: 1.00000
## 10000:14999 : 45      Max.      :3.0000  Max.      : 3.1600  Max.      : 3.56793
## (Other)      :193
##      dep4      dep5      dep6      dep7
## Min.      :0.0000  Min.      : -0.3117  Min.      : -0.1022  Min.      :0.0000
## 1st Qu.:0.0000  1st Qu.: 0.0000  1st Qu.: 0.0000  1st Qu.:0.0000
## Median :1.0000  Median : 0.0000  Median : 0.0000  Median :0.0000
## Mean      :0.7562  Mean      : 0.3564  Mean      : 0.3427  Mean      :0.3238
## 3rd Qu.:1.0000  3rd Qu.: 0.5675  3rd Qu.: 0.4545  3rd Qu.:0.0000
## Max.      :3.0000  Max.      : 3.0000  Max.      : 4.4146  Max.      :3.0000
##
##      dep8      dep9
## Min.      : -0.2742  Min.      : -0.41418
## 1st Qu.: 0.0000  1st Qu.: 0.00000
## Median : 0.0000  Median : 0.00000
## Mean      : 0.2065  Mean      : 0.06338
## 3rd Qu.: 0.0000  3rd Qu.: 0.00000
## Max.      : 3.0000  Max.      : 3.00000
##

```

## 4.2 Model creation (maybe combine with next section?)

Mostly code for now, will add more later on.

TODO:

- Use mode imputation for depression and outcome.
- Mention n rows removal in case of listwise

Summary of code:


1. Create imputed data for each ad-hoc method. 2. combine deps levels for each dataset. 3. create models for analysis.

```

##
## iter imp variable
## 1 1 dep2 dep3 dep5 dep6 dep8 dep9


```

```
## /\      /\
## { '---' }
## { 0  0 }
## ==> V <== No need for mice. This data set is completely observed.
## \  \|\ / /
##  '-----'
```



```
##      id drink_regularly sex age ethnicity education marital household_income dep
## 525  1                1  1  1          1          1          1          1  1
##      0                0  0  0          0          0          0          0  0
##
## 525  0
##      0
```

```
## /\      /\
## { '---' }
## { 0  0 }
## ==> V <== No need for mice. This data set is completely observed.
## \  \|\ / /
##  '-----'
```



```
##      id drink_regularly sex age ethnicity education marital household_income dep
## 525  1                1  1  1          1          1          1                1  1
##      0                0  0  0          0          0          0                0  0
##
## 525  0
##      0
```

### 4.3 Comparison of the two different models in terms of missing data treatment !!! (Ruben)

### 4.4 Conclusion in terms of answering RQ (Nisse)

## 5 References

- Garnett, Claire, Sabrina Kastaun, Jamie Brown, and Daniel Kotz. 2022. "Alcohol Consumption and Associations with Sociodemographic and Health-Related Characteristics in Germany: A Population Survey." *Addictive Behaviors* 125 (February): 107159. <https://doi.org/10.1016/j.addbeh.2021.107159>.
- Moore, Alison A., Robert Gould, David B. Reuben, Gail A. Greendale, M. Kallin Carter, Kefei Zhou, and Arun Karlamangla. 2005. "Longitudinal Patterns and Predictors of Alcohol Consumption in the United States." *American Journal of Public Health* 95 (3): 458–64. <https://doi.org/10.2105/ajph.2003.019471>.
- Wu, Wei, Fan Jia, and Craig Enders. 2015. "A Comparison of Imputation Strategies for Ordinal Missing Data on Likert Scale Variables." *Multivariate Behavioral Research* 50 (5): 484–503. <https://doi.org/10.1080/00273171.2015.1022644>.

## A Appendix

Table 4: Summary statistics

Variable	N	Mean	Std. Dev.	Min	Pctl. 25	Pctl. 75	Max
id	525	46470	2898	41531	43912	48934	51610
drink_regularly	446						
... yes	307	69%					
... no	139	31%					
sex	525						
... male	254	48%					
... female	271	52%					
age	525	45	14	20	33	57	69
ethnicity	525						
... mexican_american	95	18%					
... other_hispanic	61	12%					
... non-hispanic_white	220	42%					
... non-hispanic_black	124	24%					
... other	25	5%					
education	525						
... no_high_school	58	11%					
... some_high_school	101	19%					
... high_school_grad	123	23%					
... some_college	155	30%					
... college_grad	88	17%					
marital	525						
... married	279	53%					
... widowed	19	4%					
... divorced	67	13%					
... separated	14	3%					
... never_married	102	19%					
... living_with_partner	44	8%					
household_income	525						
... 0:4999	13	2%					
... 5000:9999	24	5%					
... 10000:14999	45	9%					
... 15000:19999	40	8%					
... 20000:24999	52	10%					
... 25000:34999	59	11%					
... 35000:44999	51	10%					
... 45000:54999	44	8%					
... 55000:64999	35	7%					
... 65000:74999	37	7%					
... 75000:99999	49	9%					
... 100000+	76	14%					
dep1	525	0.41	0.79	0	0	1	3
dep2	394	0.28	0.58	0	0	0	3
dep3	394	0.53	0.86	0	0	1	3
dep4	525	0.76	0.9	0	0	1	3
dep5	394	0.31	0.7	0	0	0	3
dep6	394	0.2	0.56	0	0	0	3
dep7	525	0.32	0.71	0	0	0	3
dep8	473	0.2	0.59	0	0	0	3
dep9	449	0.067	0.376	0	0	0	3